



به نام خدا

## بخش چهاردهم

# Transformers

# حمیدرضا برادران کاشانی



# Neural Machine Translation

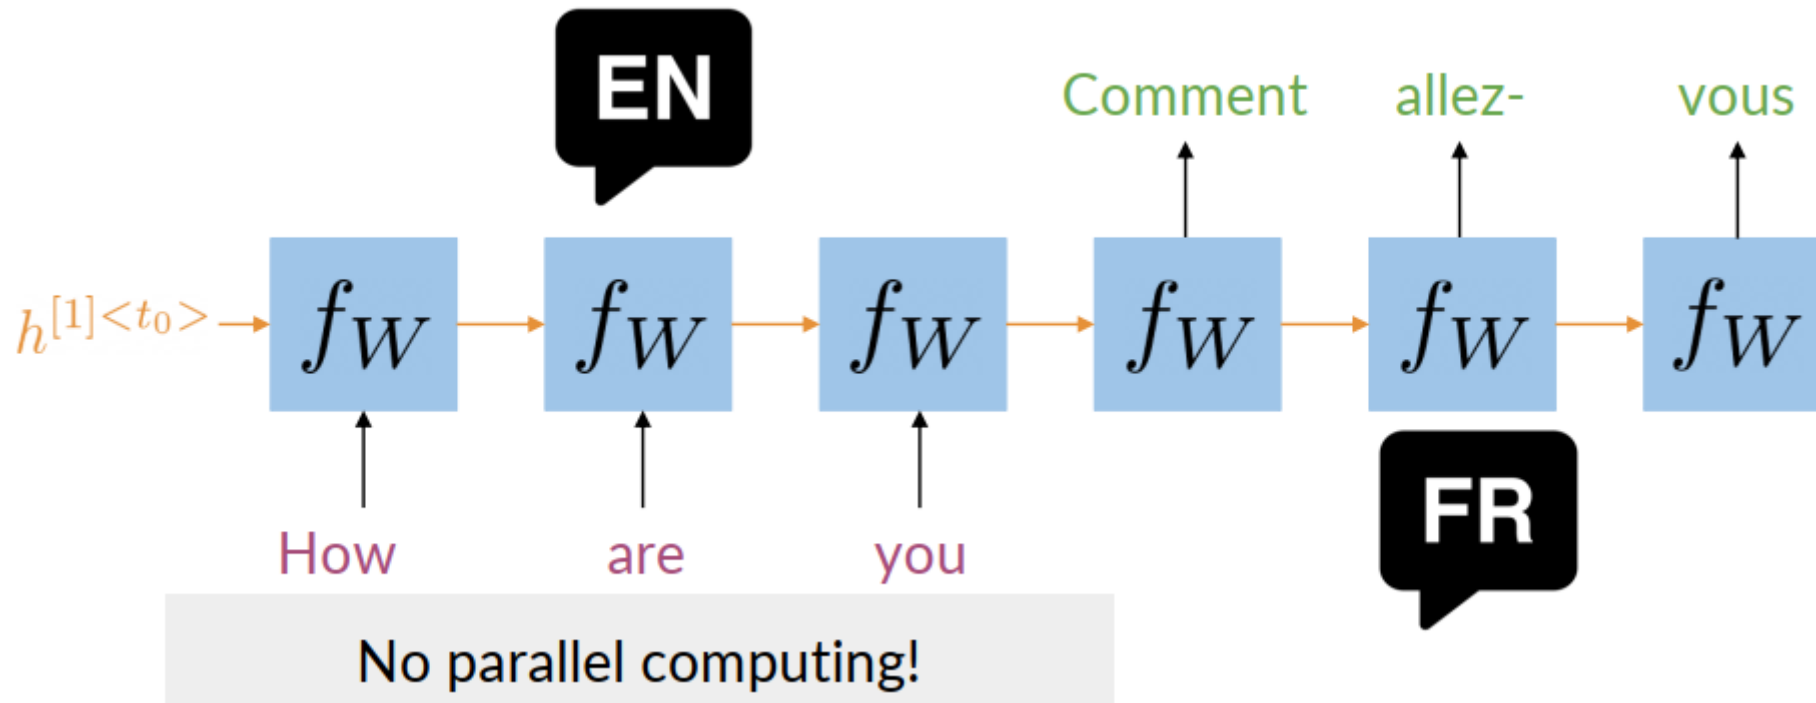


Figure from deeplearning.ai

# Sequence-to-sequence architectures

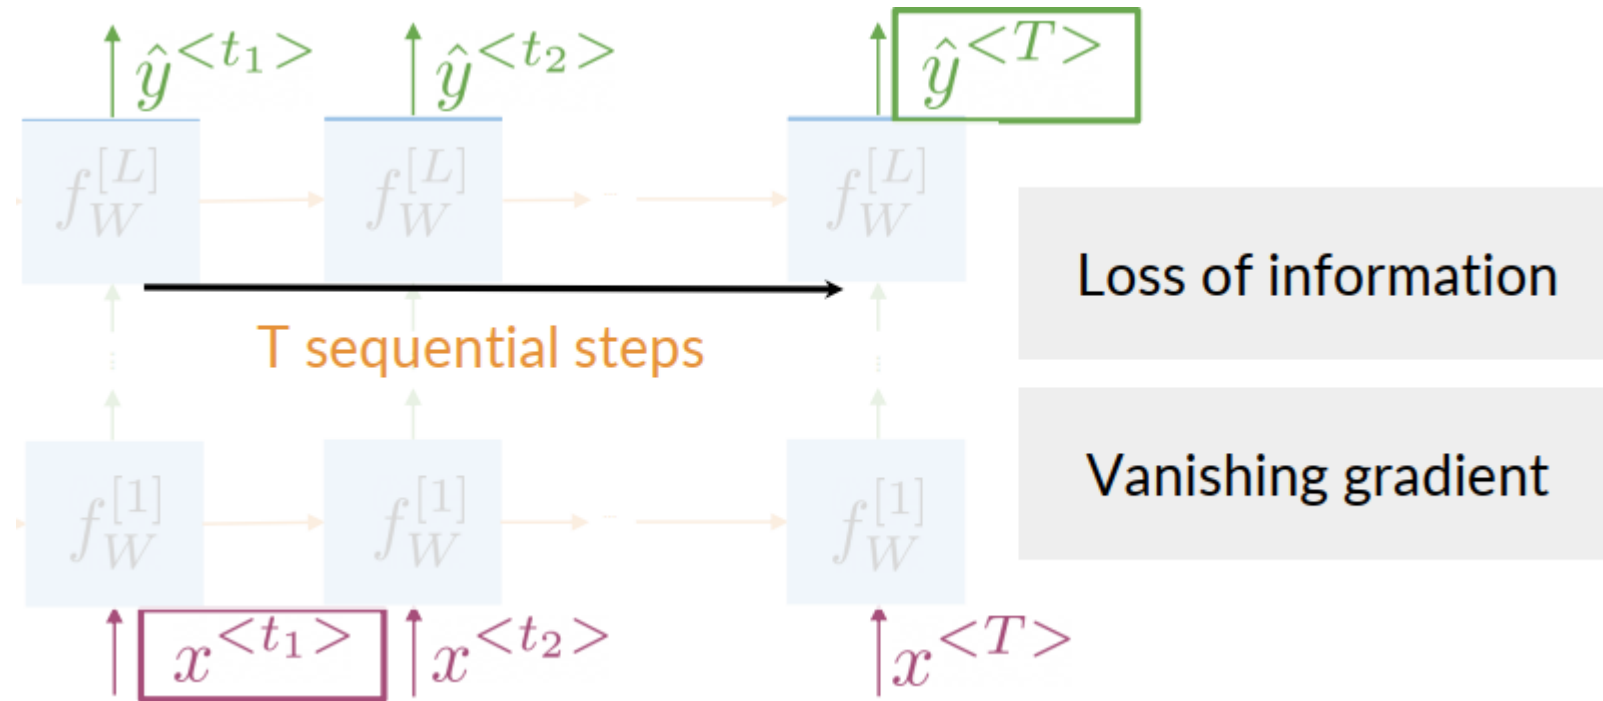
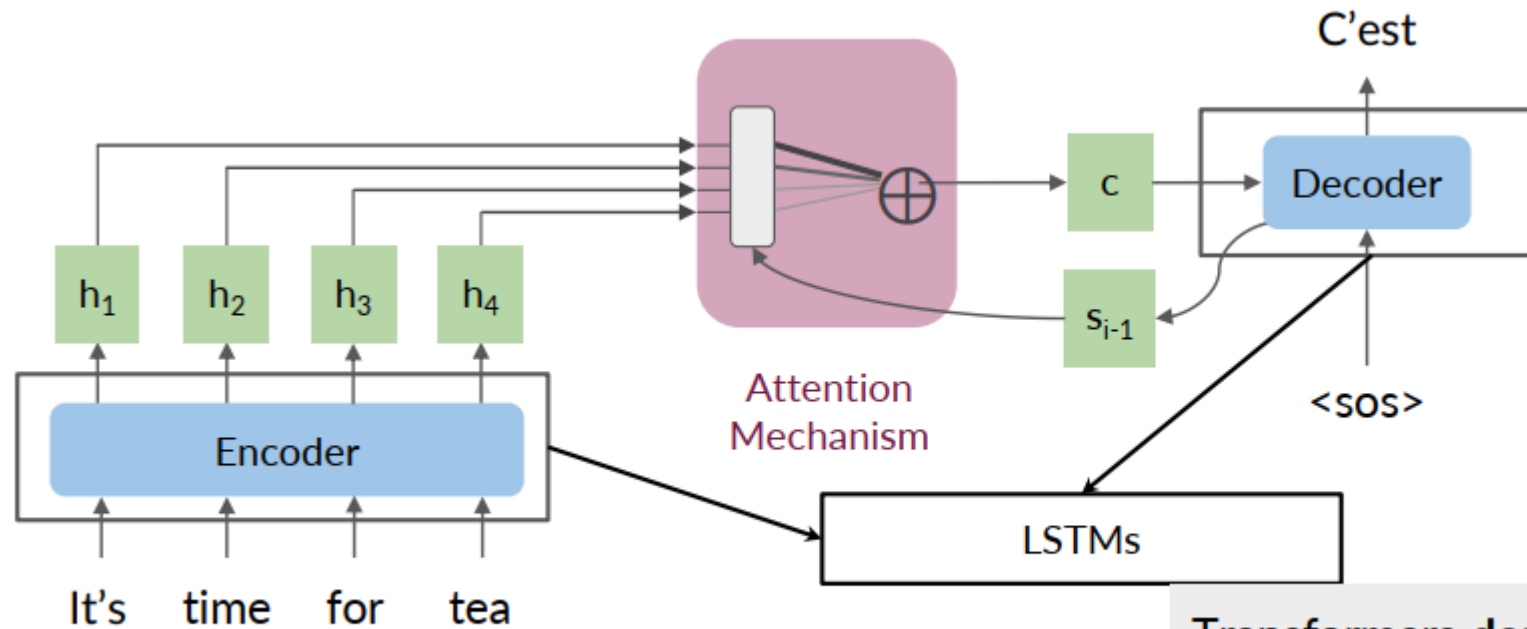


Figure from deeplearning.ai



# RNN vs. Transformers



Transformers **don't** use RNNs, such as LSTMs or GRUs



Figure from deeplearning.ai



# The Transformer Model

---

## Attention Is All You Need

---

**Ashish Vaswani\***  
Google Brain  
avaswani@google.com

**Noam Shazeer\***  
Google Brain  
noam@google.com

**Niki Parmar\***  
Google Research  
nikip@google.com

**Jakob Uszkoreit\***  
Google Research  
usz@google.com

**Llion Jones\***  
Google Research  
llion@google.com

**Aidan N. Gomez\* †**  
University of Toronto  
aidan@cs.toronto.edu

**Łukasz Kaiser\***  
Google Brain  
lukaszkaiser@google.com

**Illia Polosukhin\* ‡**  
illia.polosukhin@gmail.com


<https://arxiv.org/abs/1706.03762>



Figure from deeplearning.ai



# Scaled Dot-Product Attention



(Vaswani et al., 2017)

$$\text{softmax} \left( \frac{QK^{\top}}{\sqrt{d_k}} \right) V$$



Figure from deeplearning.ai



# Multi-Head Attention

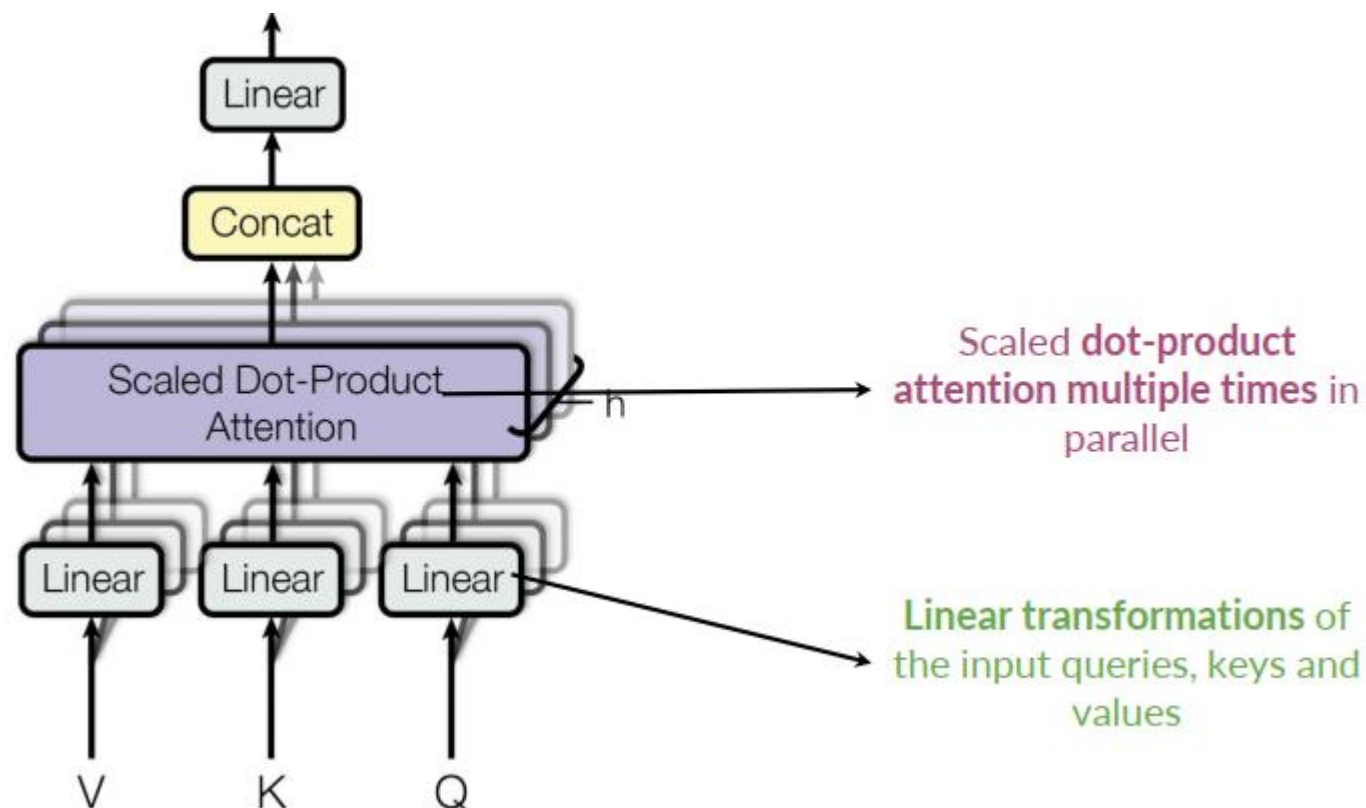


Figure from deeplearning.ai

# The Encoder

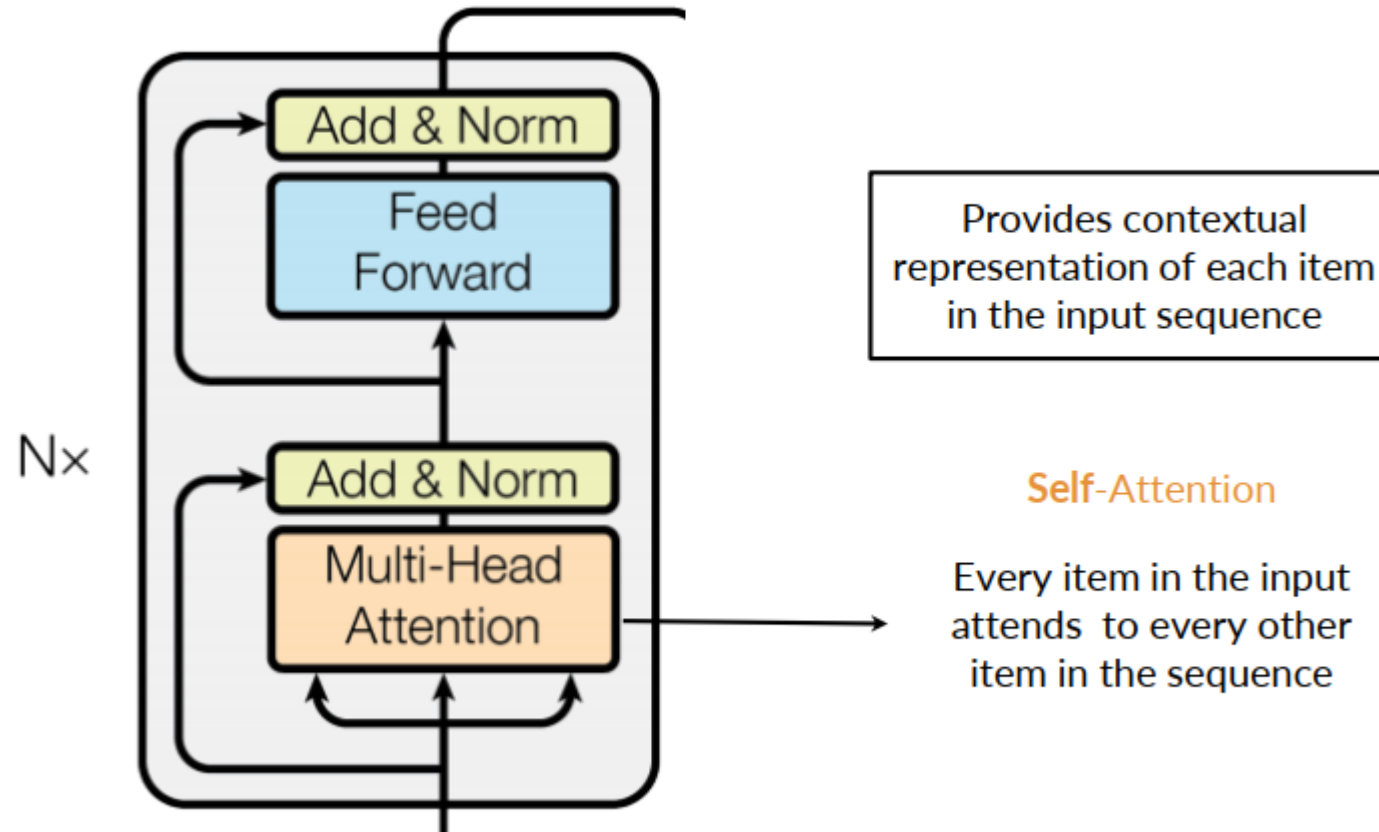


Figure from deeplearning.ai





# The Decoder

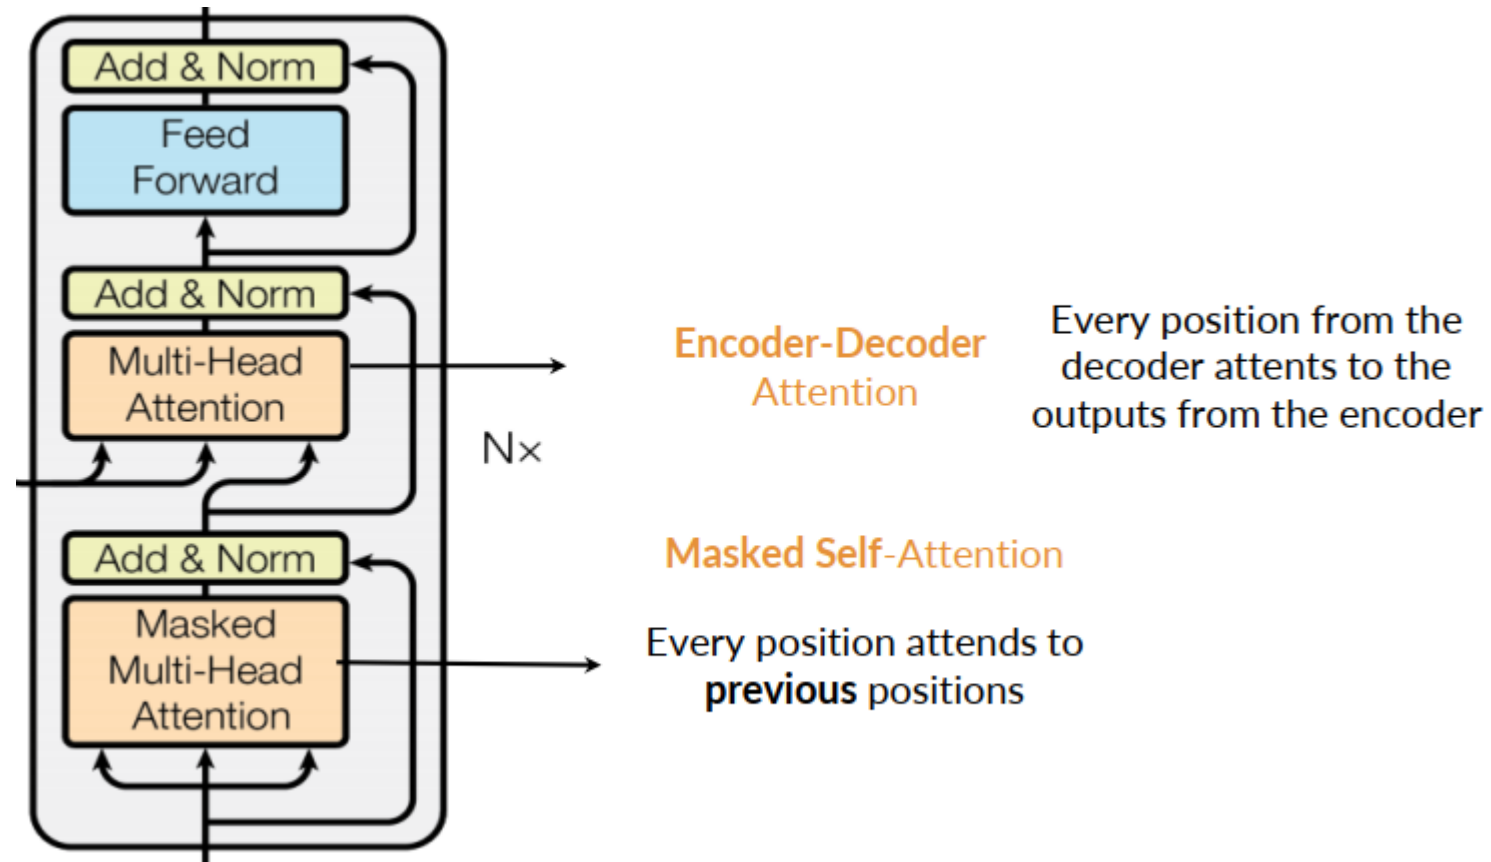


Figure from deeplearning.ai



# RNNs vs Transformer: Positional Encoding

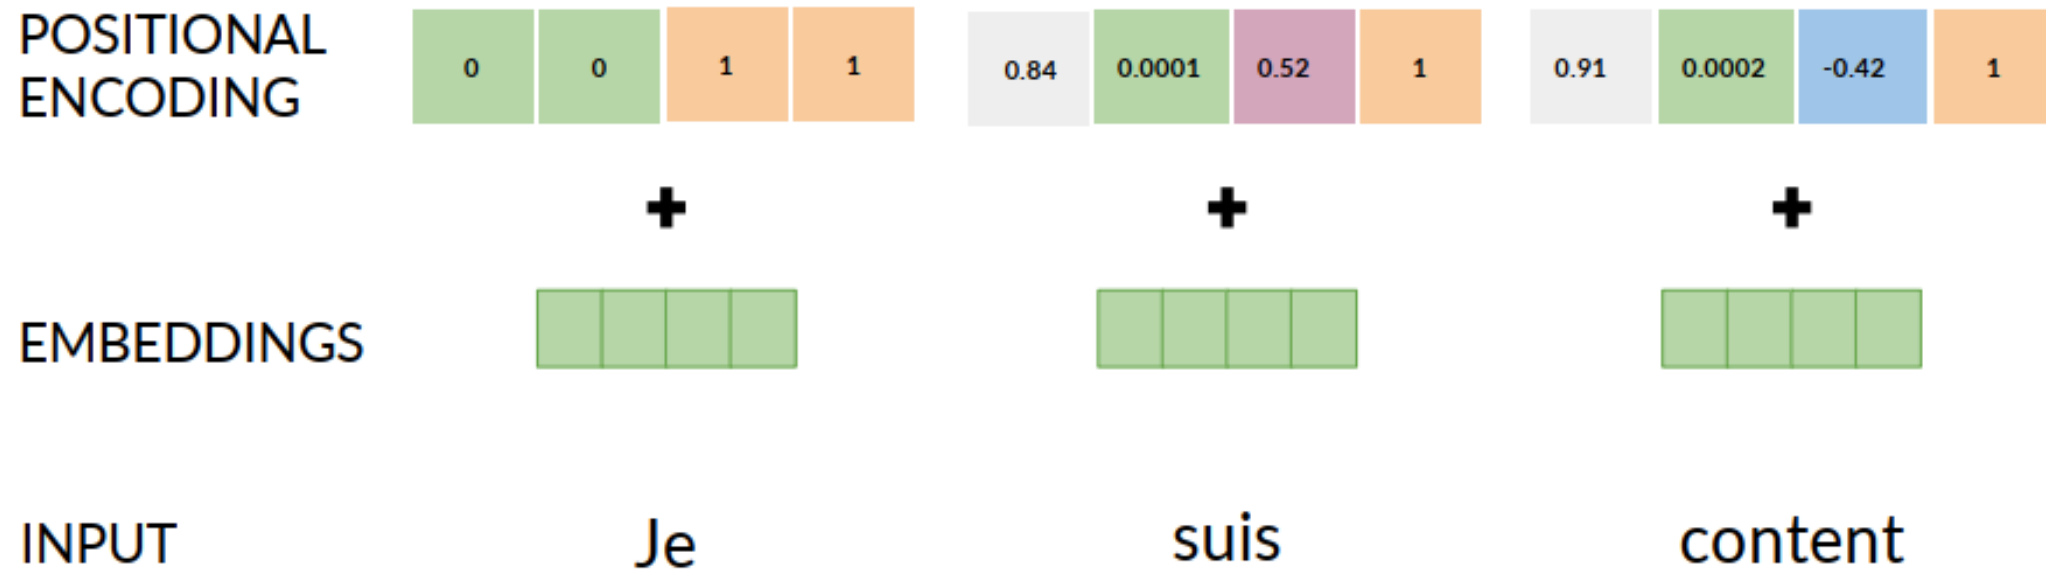


Figure from deeplearning.ai



# The Transformer Model

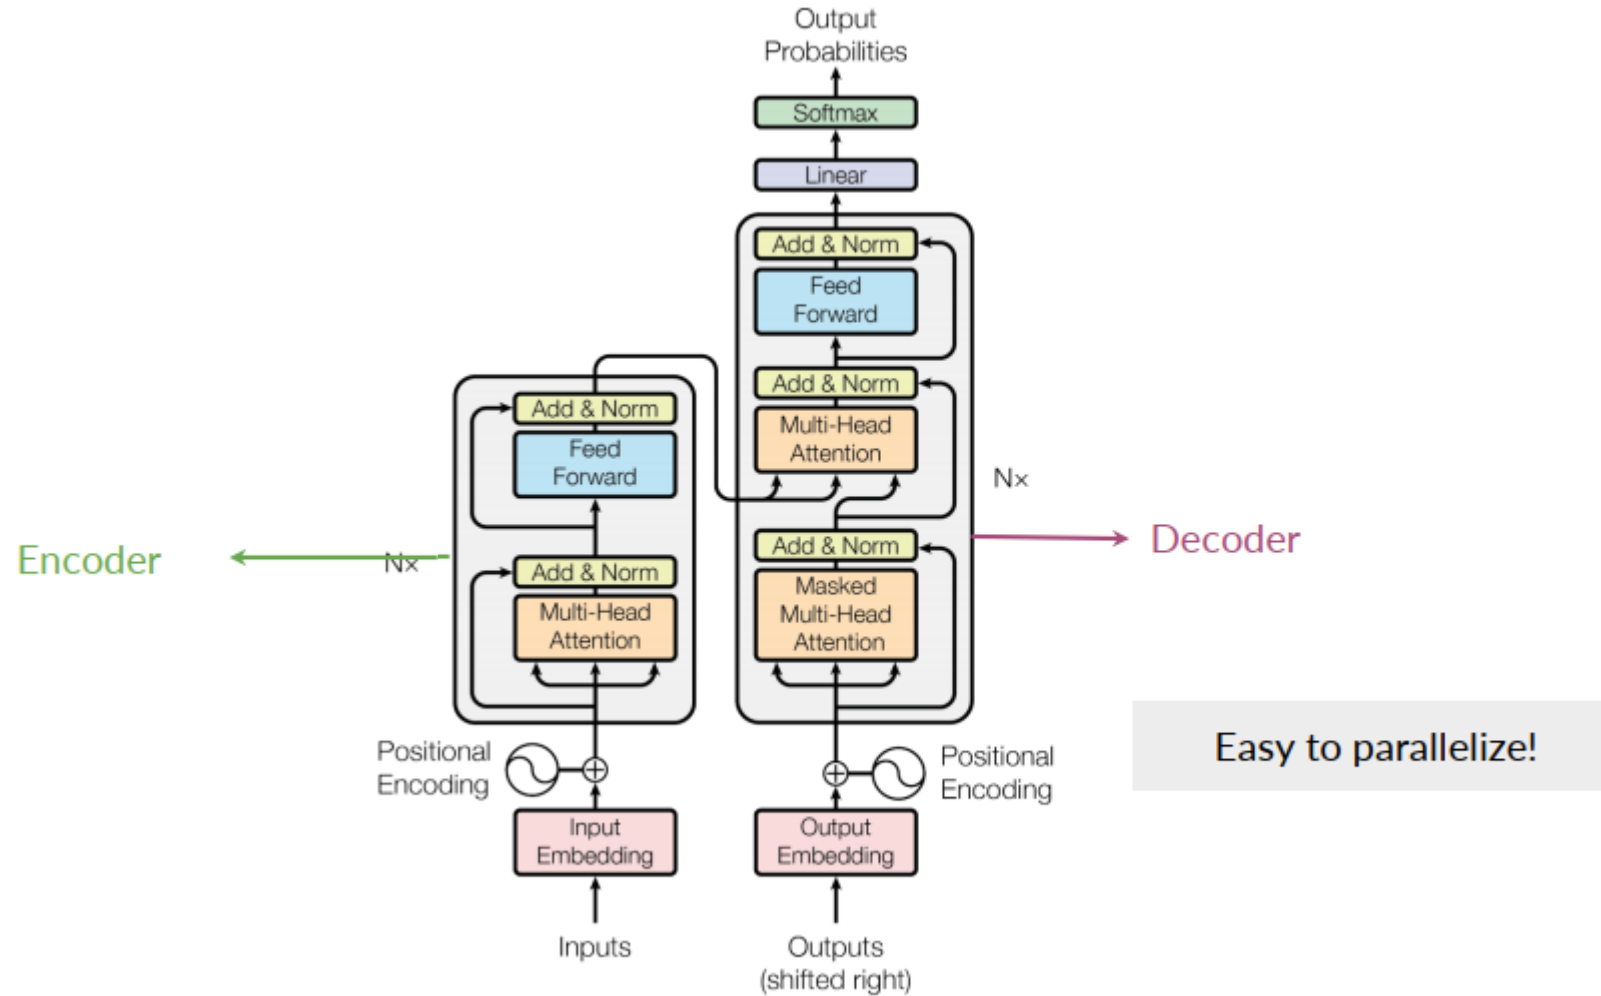


Figure from deeplearning.ai

# Summary - 1

---



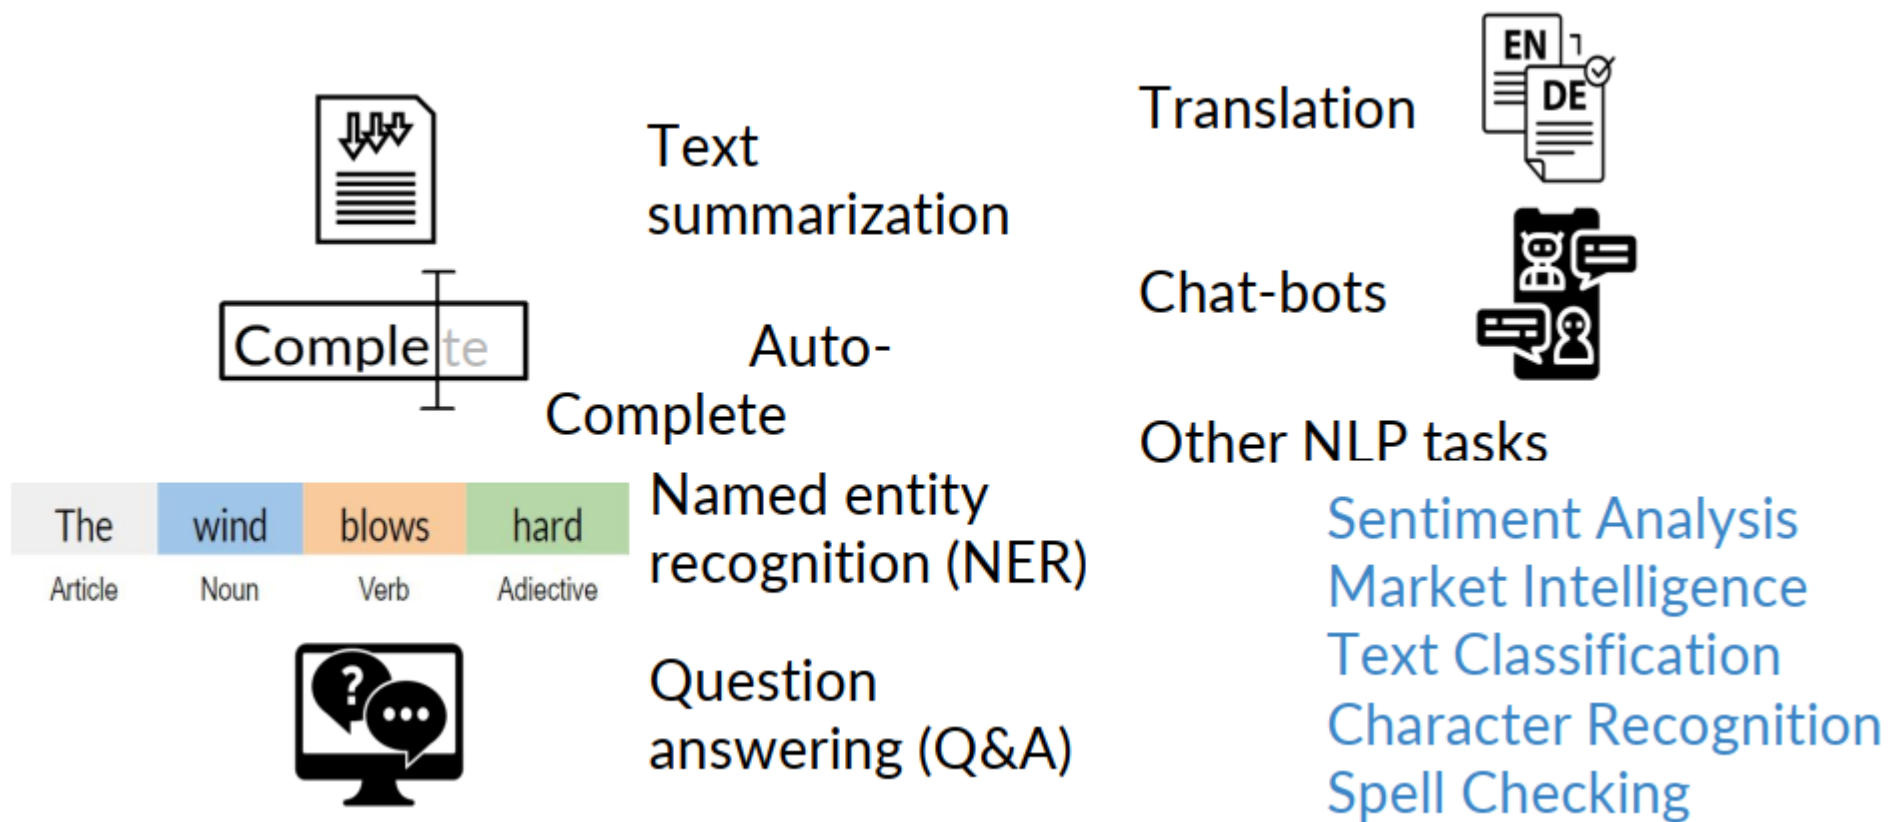
- In RNNs parallel computing is difficult to implement
- For long sequences in RNNs there is loss of information
- In RNNs there is the problem of vanishing gradient
- Transformers help with all of the above



Figure from deeplearning.ai



# Transformer NLP applications





# State of the Art Transformers

---

Radford, A., et al. (2018)  
Open AI

Devlin, J., et al. (2018)  
Google AI Language

Colin, R., et al. (2019)  
Google

GPT-2: Generative Pre-training for  
Transformer

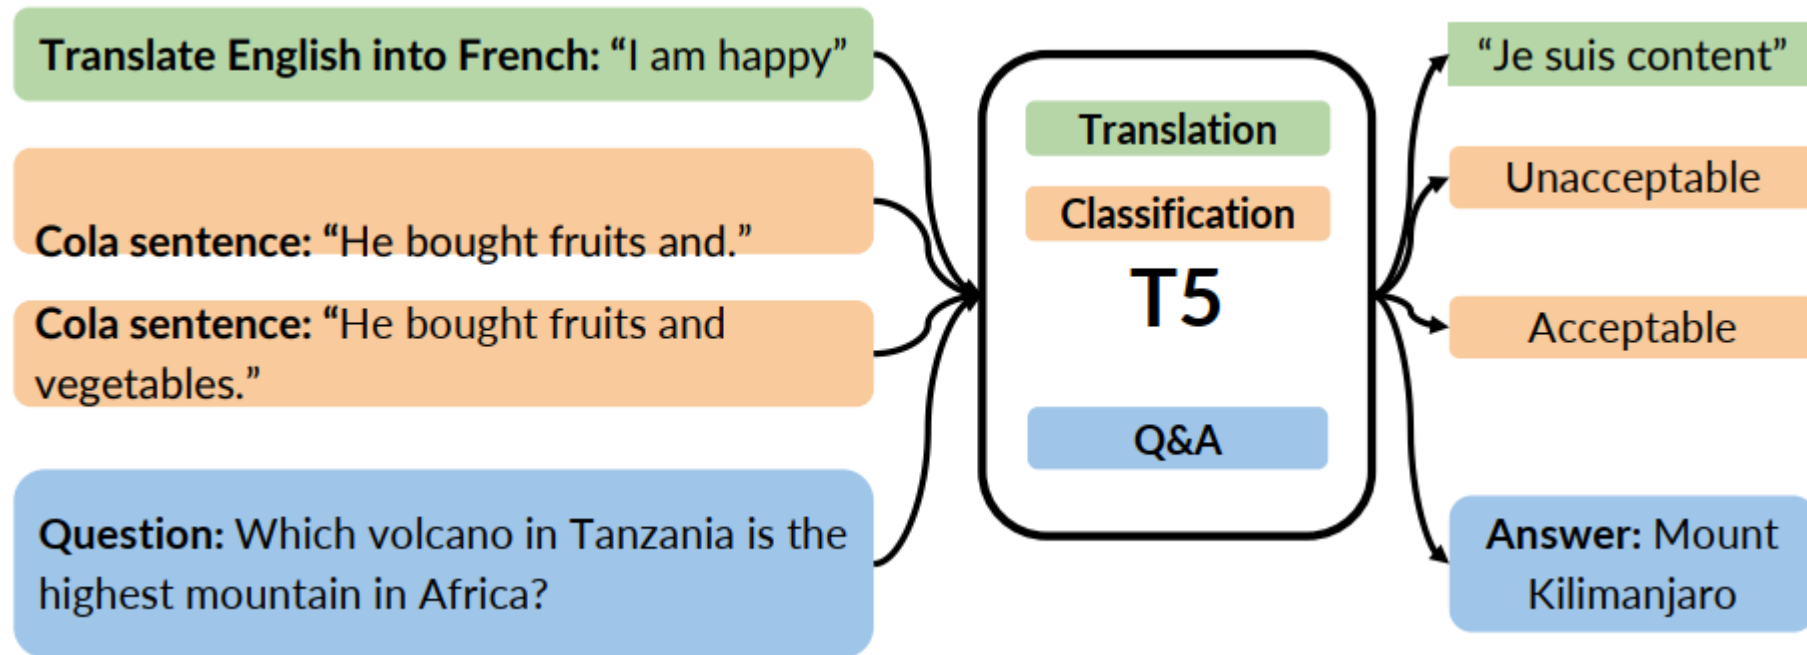
BERT: Bidirectional Encoder  
Representations from Transformers

T5: Text-to-text transfer transformer



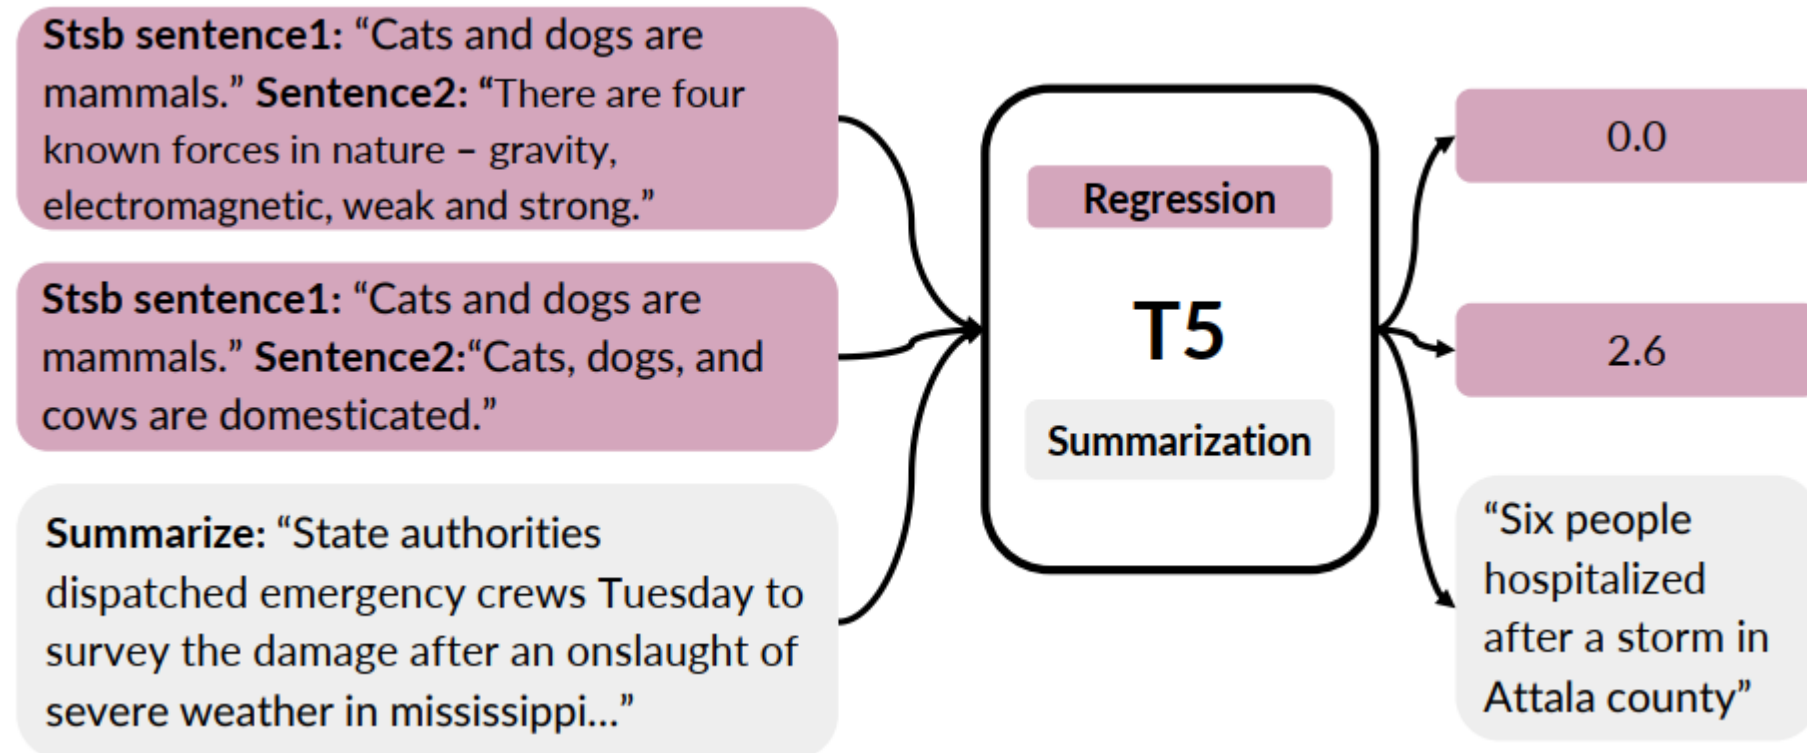


# T5: Text-To-Text Transfer Transformer






# T5: Text-To-Text Transfer Transformer







# Scaled dot-product attention

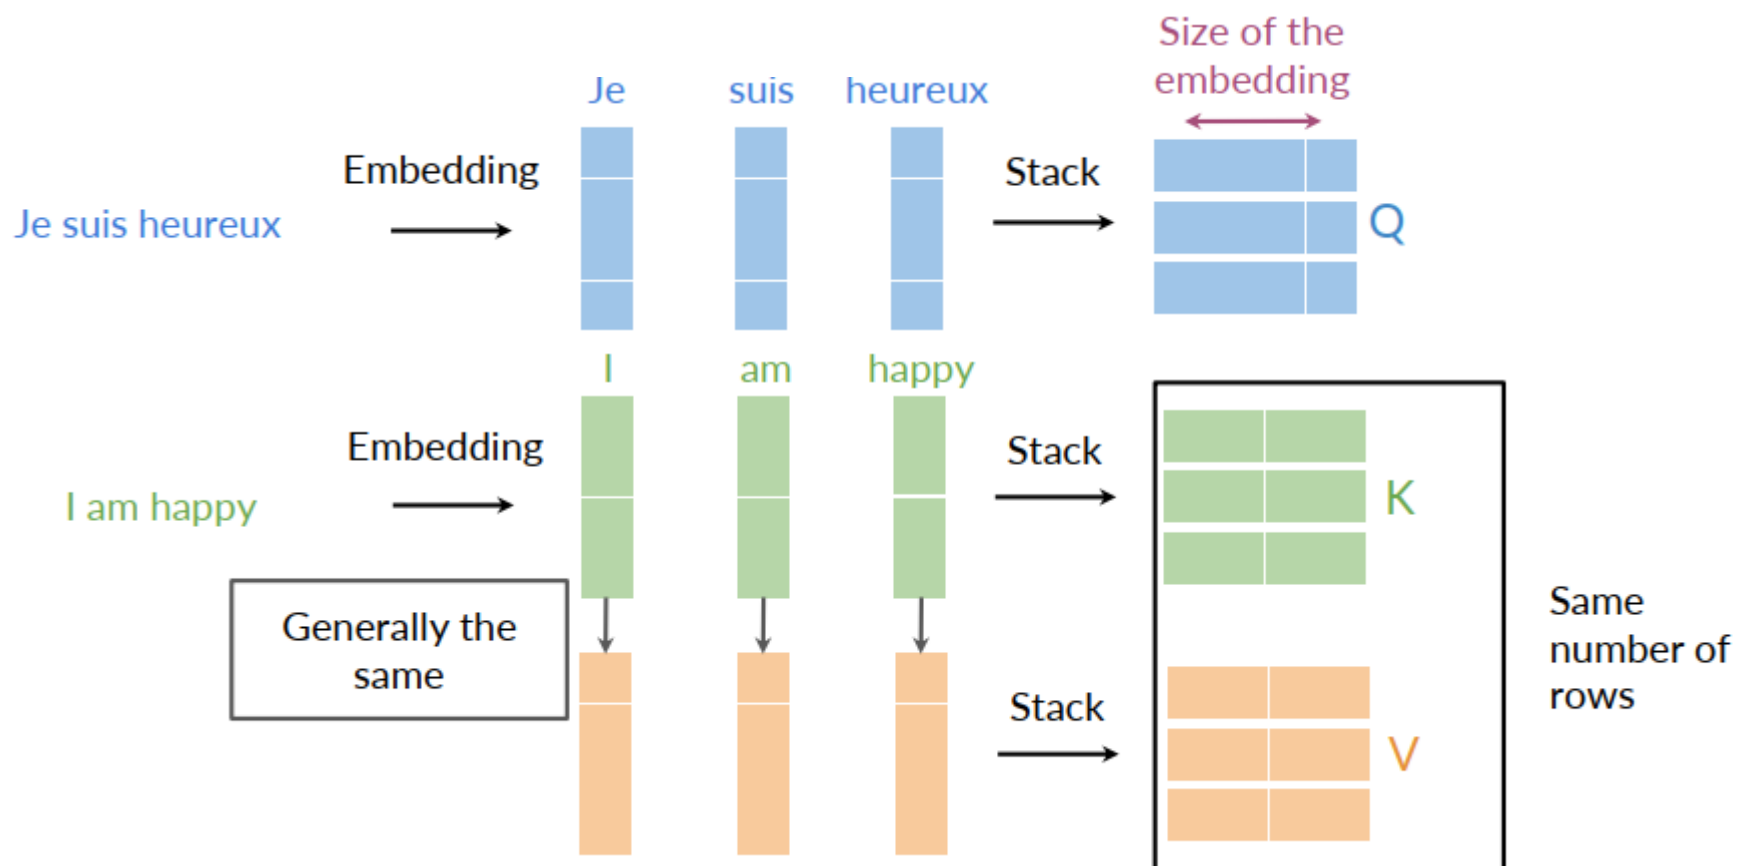


Just two matrix multiplications  
and a Softmax!







# Queries, Keys and Values





# Attention Math



$$\text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

Figure from deeplearning.ai



# Three ways of attention

## 3.2.3 Applications of Attention in our Model

The Transformer uses multi-head attention in three different ways:

- In "encoder-decoder attention" layers, the queries come from the previous decoder layer, and the memory keys and values come from the output of the encoder. This allows every position in the decoder to attend over all positions in the input sequence. This mimics the typical encoder-decoder attention mechanisms in sequence-to-sequence models such as [38, 2, 9].
- The encoder contains self-attention layers. In a self-attention layer all of the keys, values and queries come from the same place, in this case, the output of the previous layer in the encoder. Each position in the encoder can attend to all positions in the previous layer of the encoder.
- Similarly, self-attention layers in the decoder allow each position in the decoder to attend to all positions in the decoder up to and including that position. We need to prevent leftward information flow in the decoder to preserve the auto-regressive property. We implement this inside of scaled dot-product attention by masking out (setting to  $-\infty$ ) all values in the input of the softmax which correspond to illegal connections. See Figure 2.





# Queries, Keys, values and Attention



Figure from deeplearning.ai



# Encoder-Decoder Attention

Queries from one sentence, **keys** and **values** from another




Figure from deeplearning.ai



# Self-Attention

Queries, keys and values come from the **same sentence**




Figure from deeplearning.ai



# Masked Self-Attention

Queries, keys and values come from the **same sentence**. Queries don't attend to future positions.



Figure from deeplearning.ai






# Masked self-attention math

Diagram illustrating the masked self-attention mechanism:

$$\text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} + \text{Mask} \right) V$$

Legend:  → Minus infinity

The Mask matrix (3x3) is:

0		
0	0	
0	0	0

The output matrix  $V$  (3x4) is represented by orange squares.

Weights assigned to future positions are equal to 0

The mask matrix (3x3) is:

	0	0
		0



# Multi-Head Attention - Overview

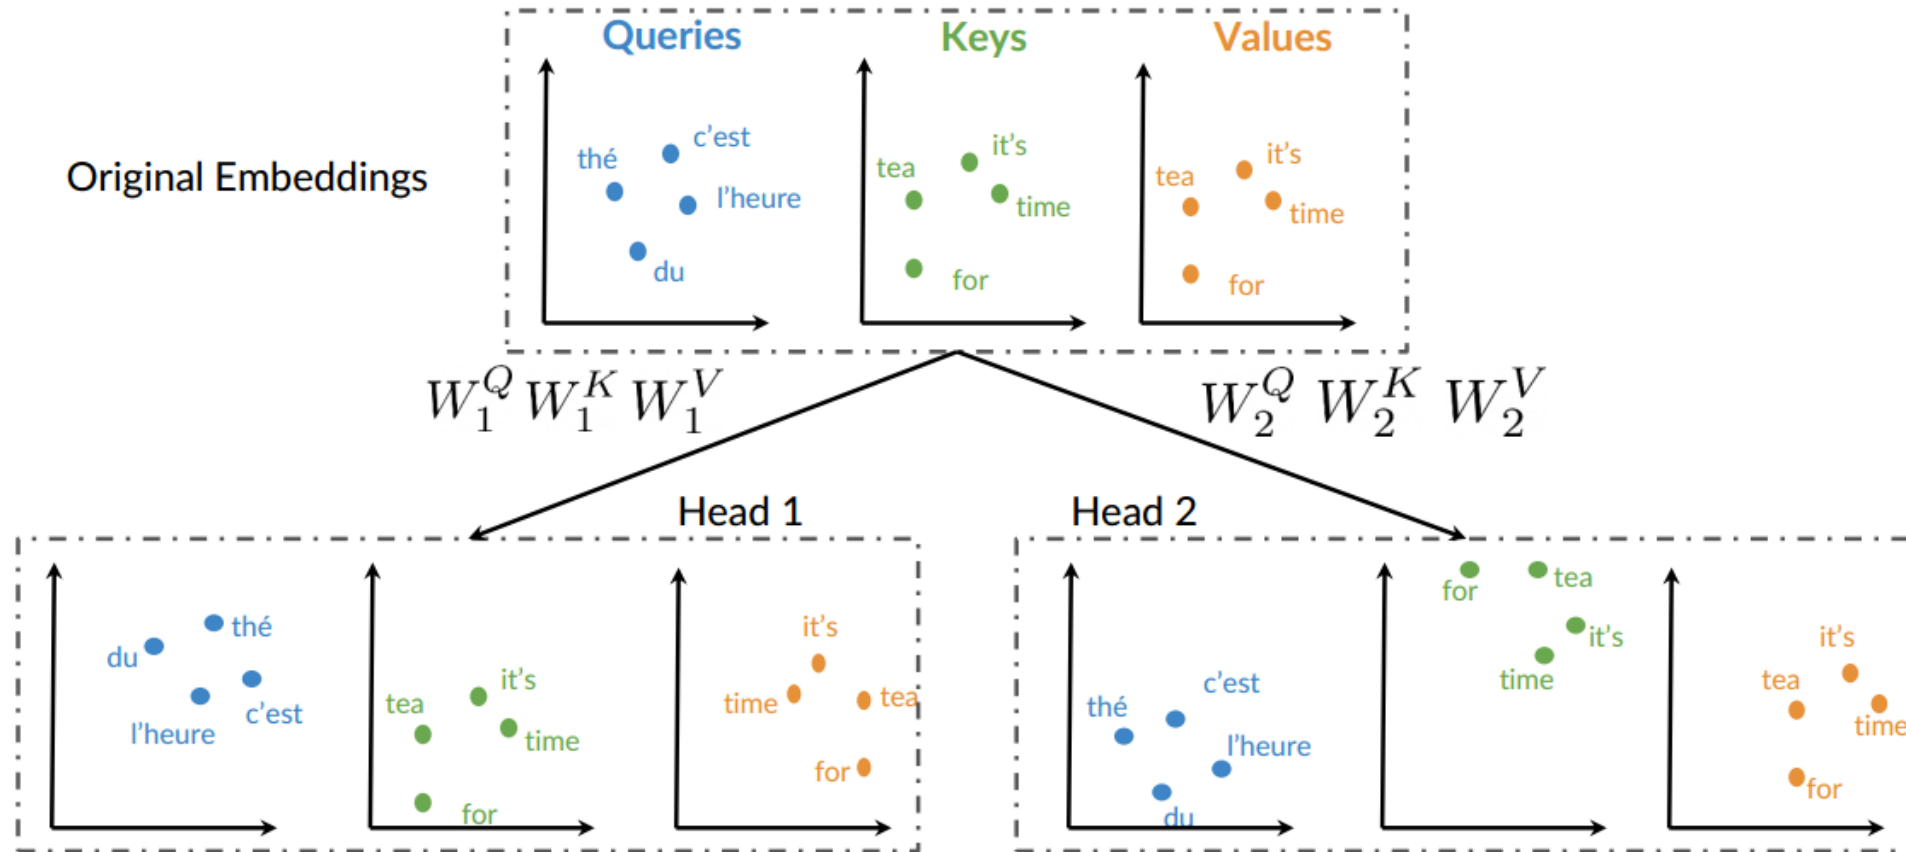


Figure from deeplearning.ai



# Multi-Head Attention - Overview

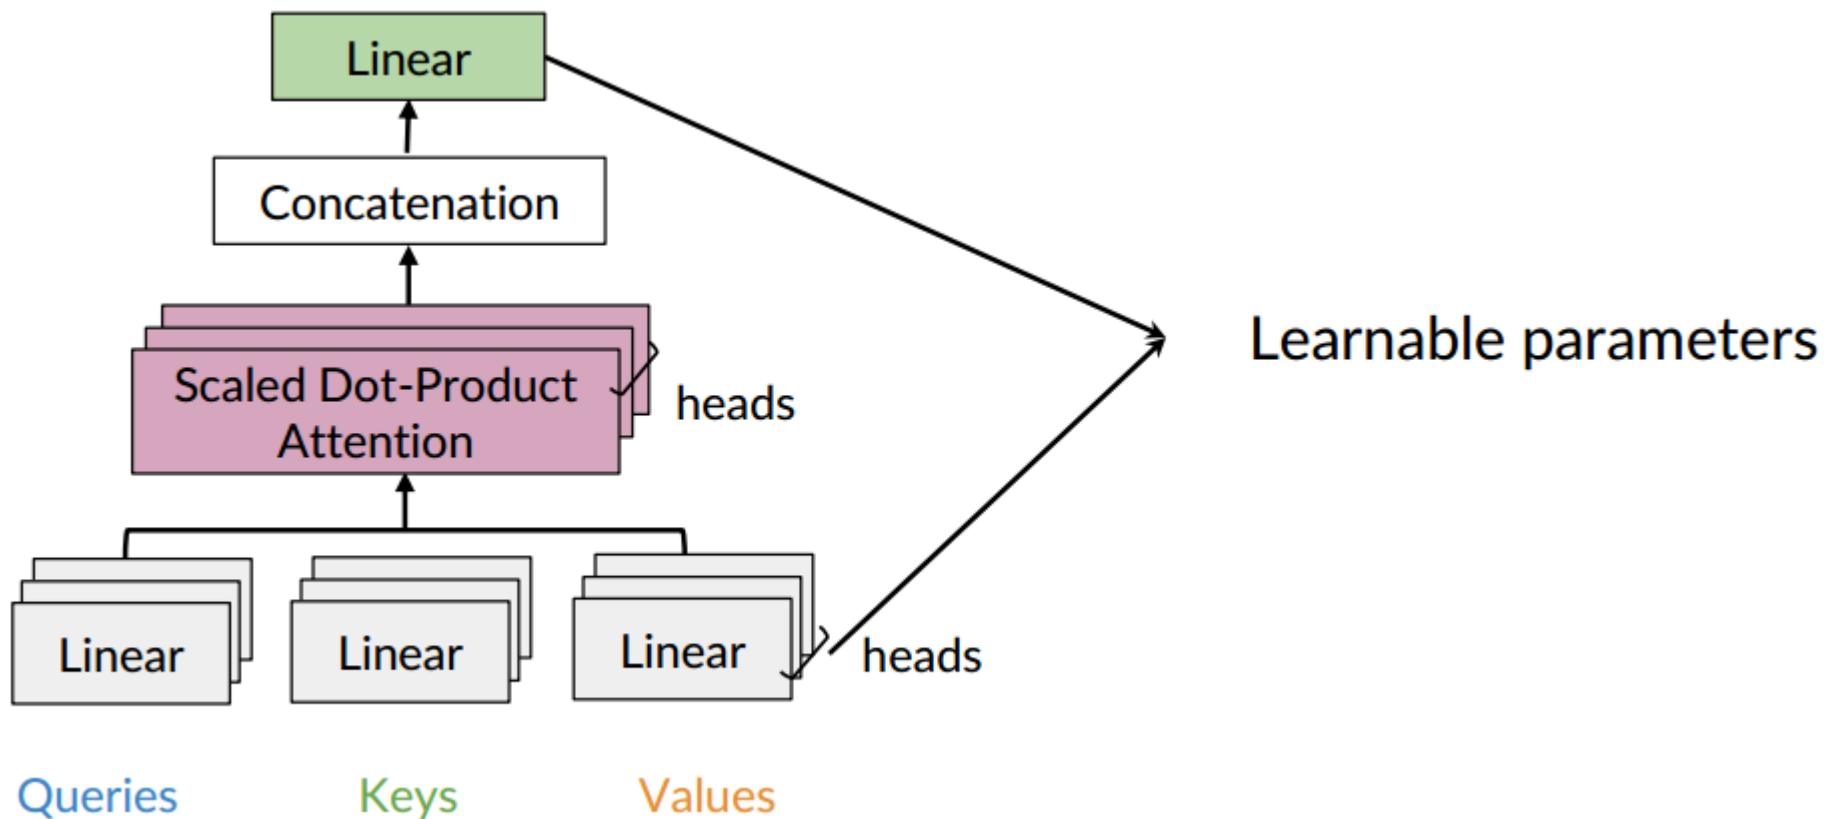
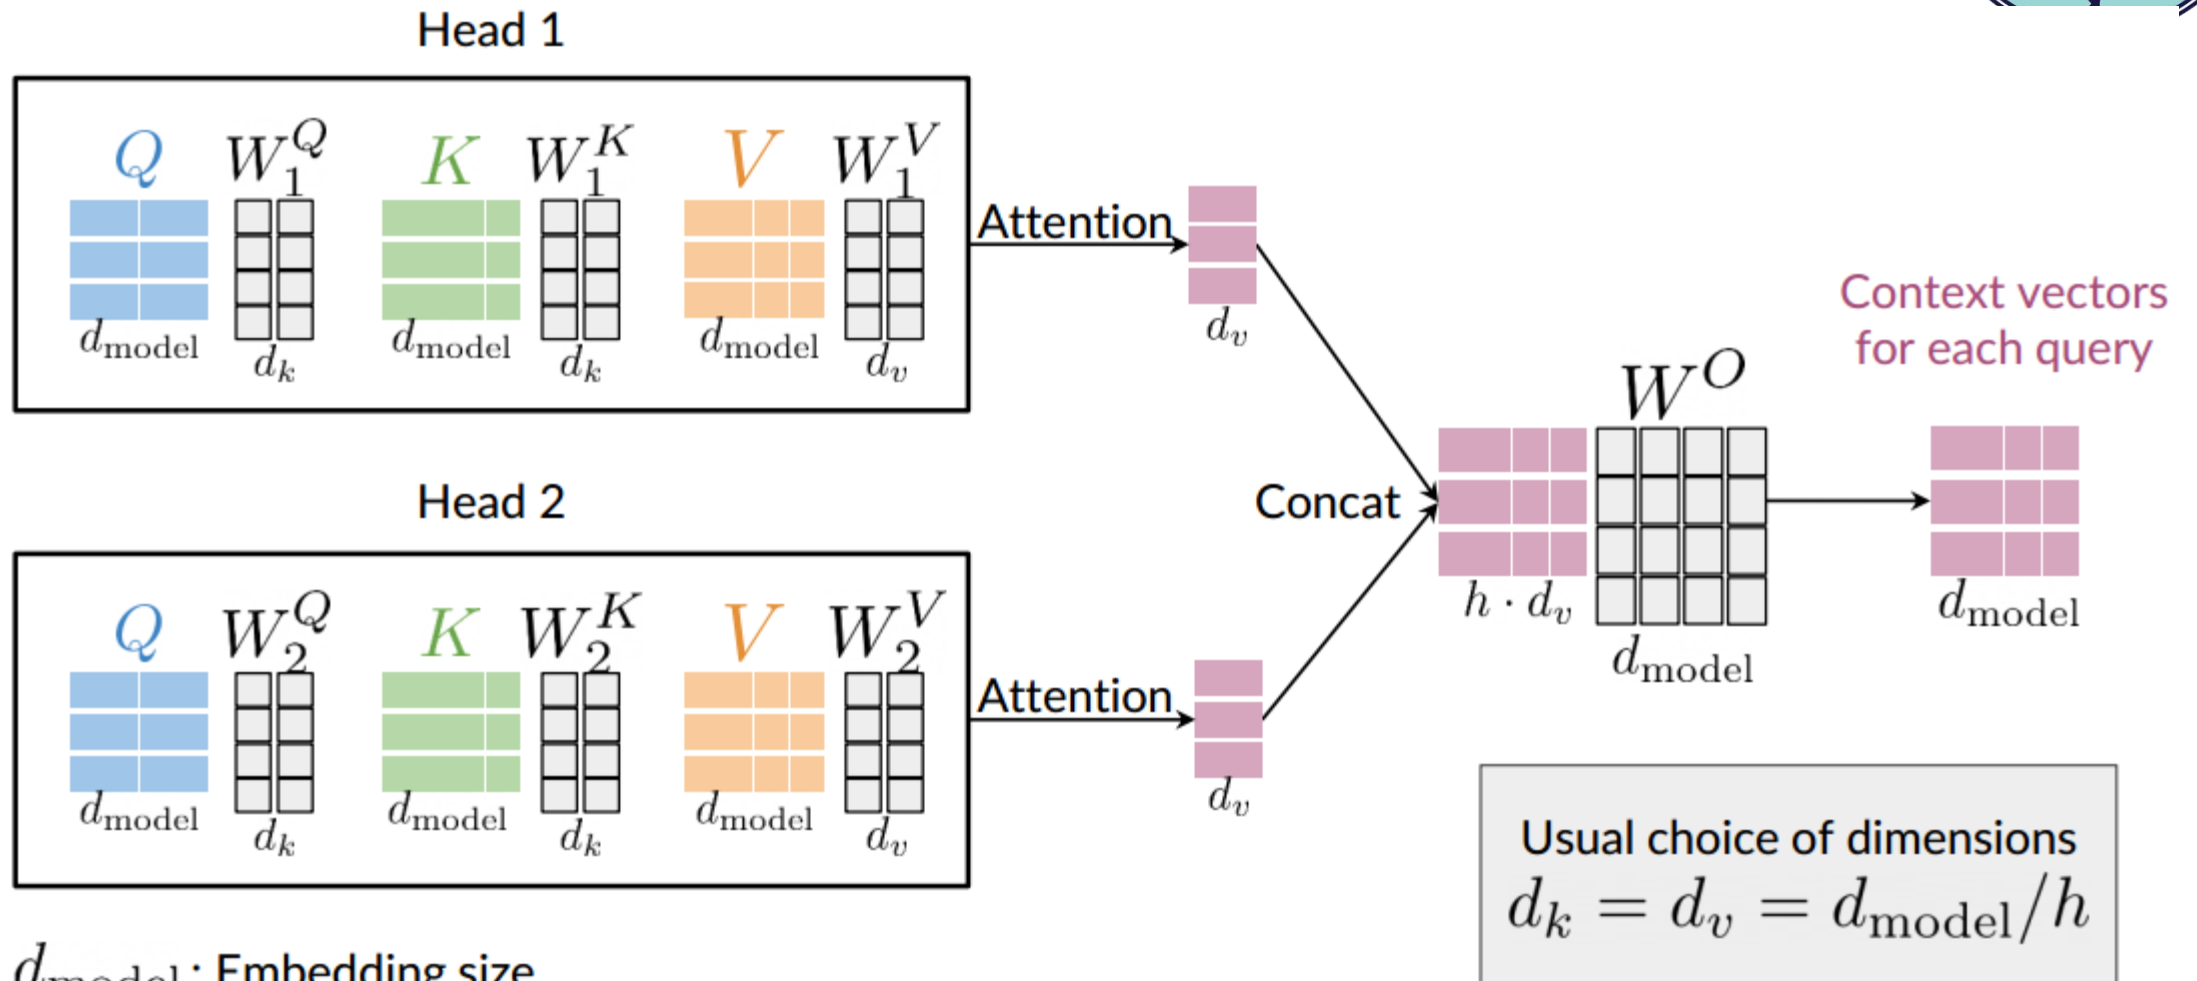


Figure from deeplearning.ai



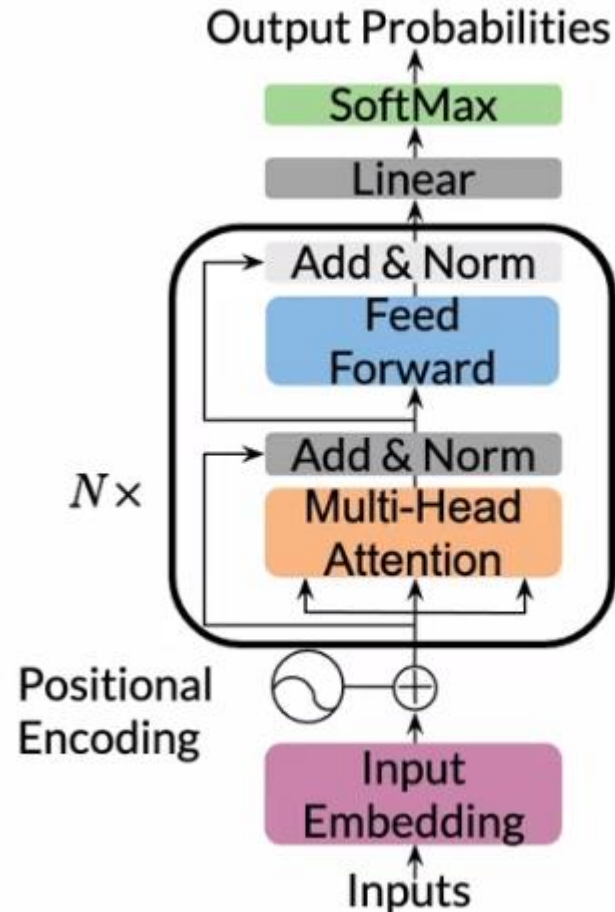
# Multi-Head Attention



$d_{\text{model}}$ : Embedding size  
Figure from deeplearning.ai



# GPT-2: Transformer Decoder



## Overview

- input: sentence or paragraph
  - we predict the next word
- sentence gets embedded, add positional encoding
  - (vectors representing  $\{0, 1, 2, \dots, K\}$ )
- multi-head attention looks at previous words
- feed-forward layer with ReLU
  - that's where most parameters are!
- residual connection with layer normalization
- repeat  $N$  times
- dense layer and softmax for output



Figure from deeplearning.ai



# GPT-2: Transformer Decoder

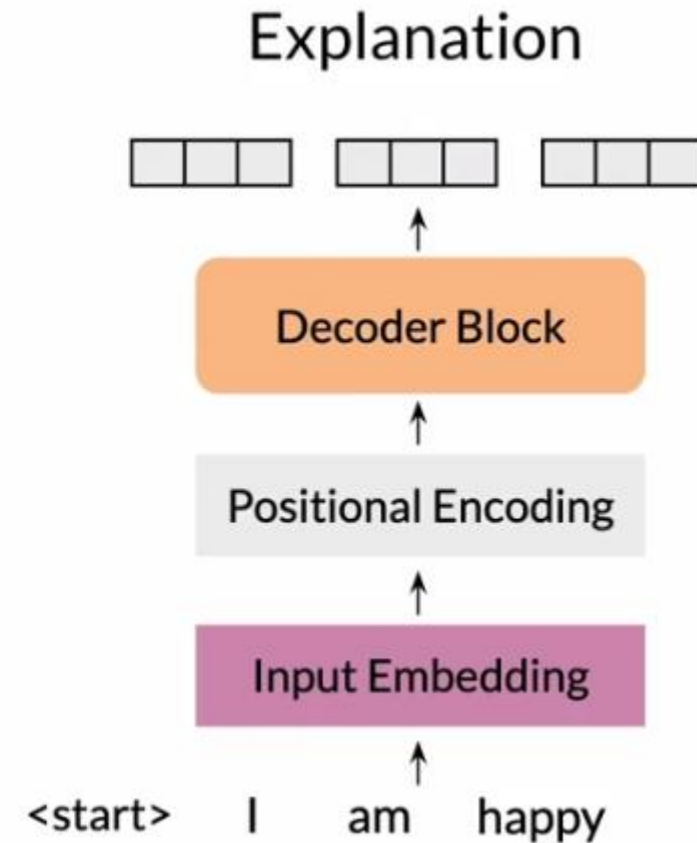



Figure from deeplearning.ai





# GPT-2: Transformer Decoder

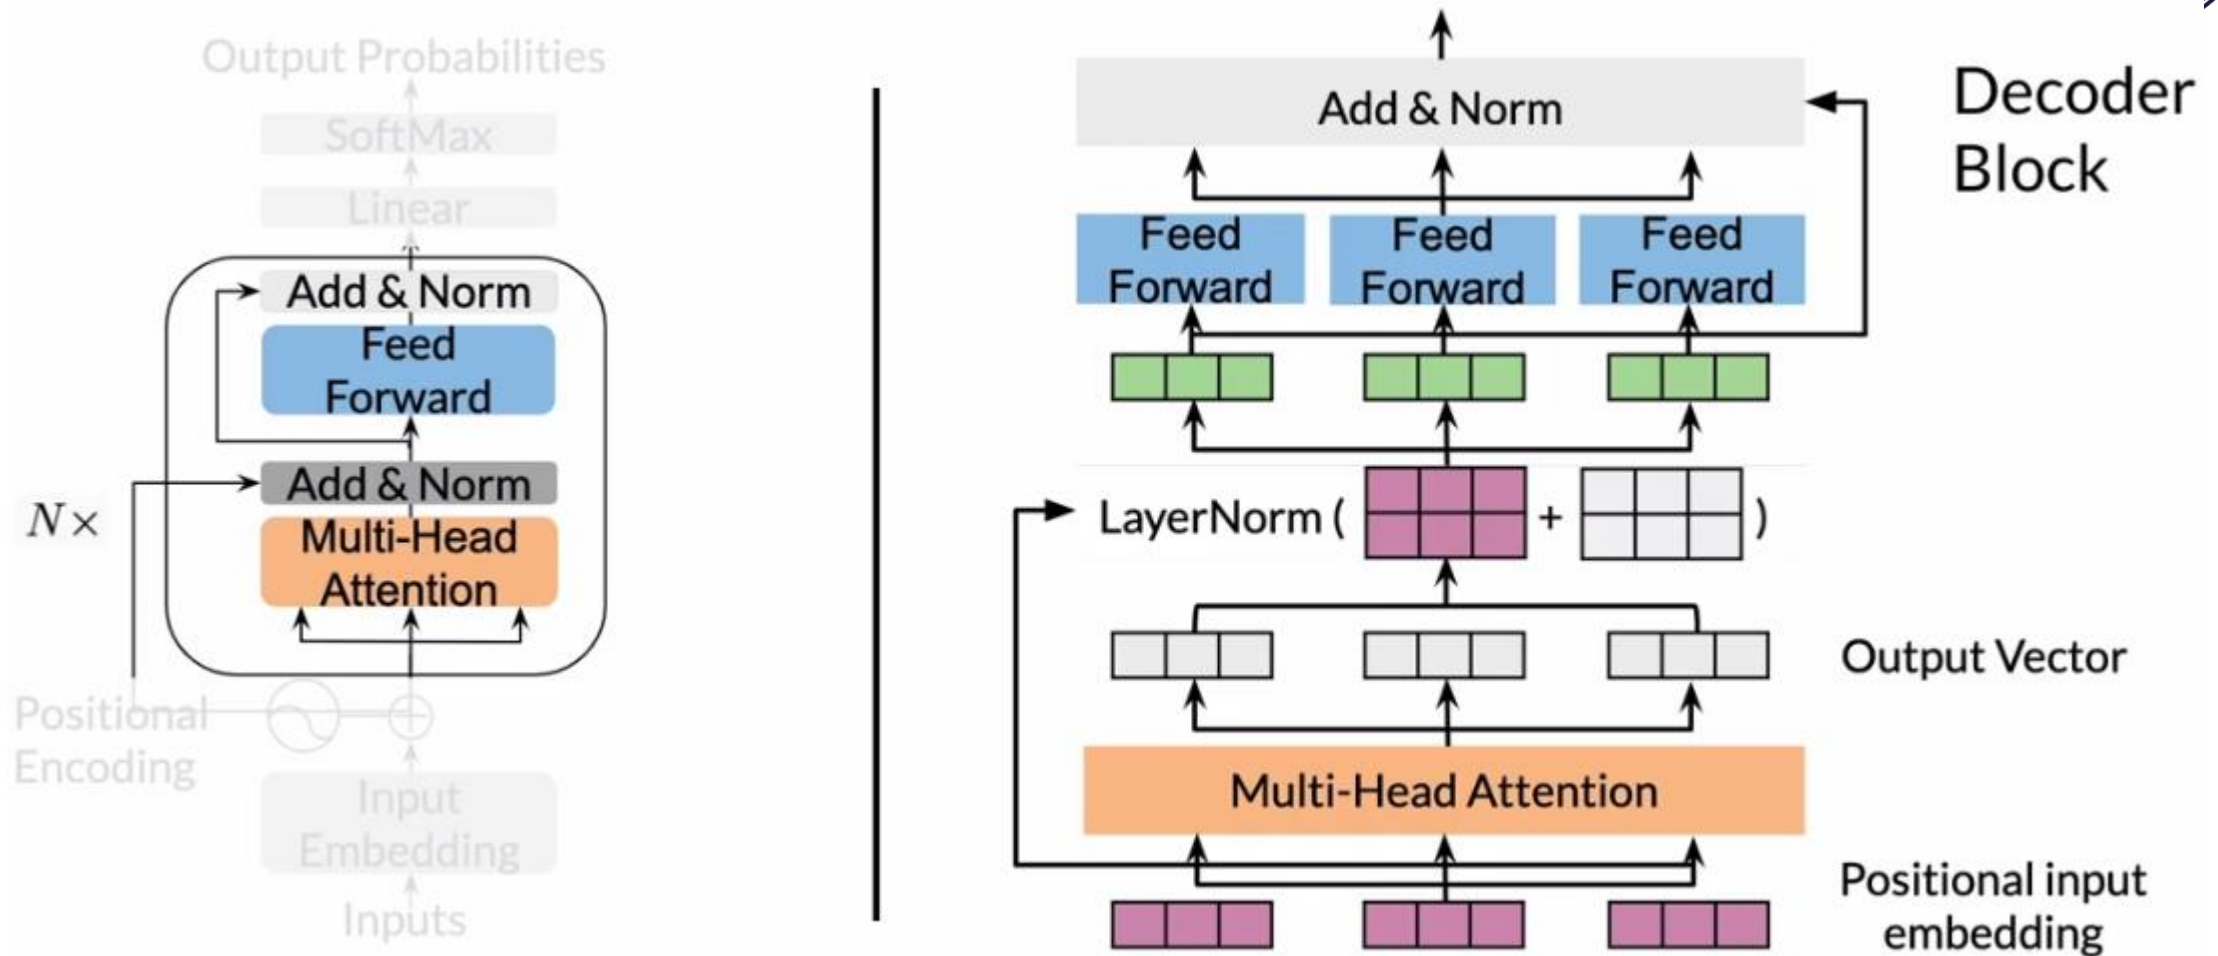


Figure from deeplearning.ai



# GPT-2: Transformer Decoder

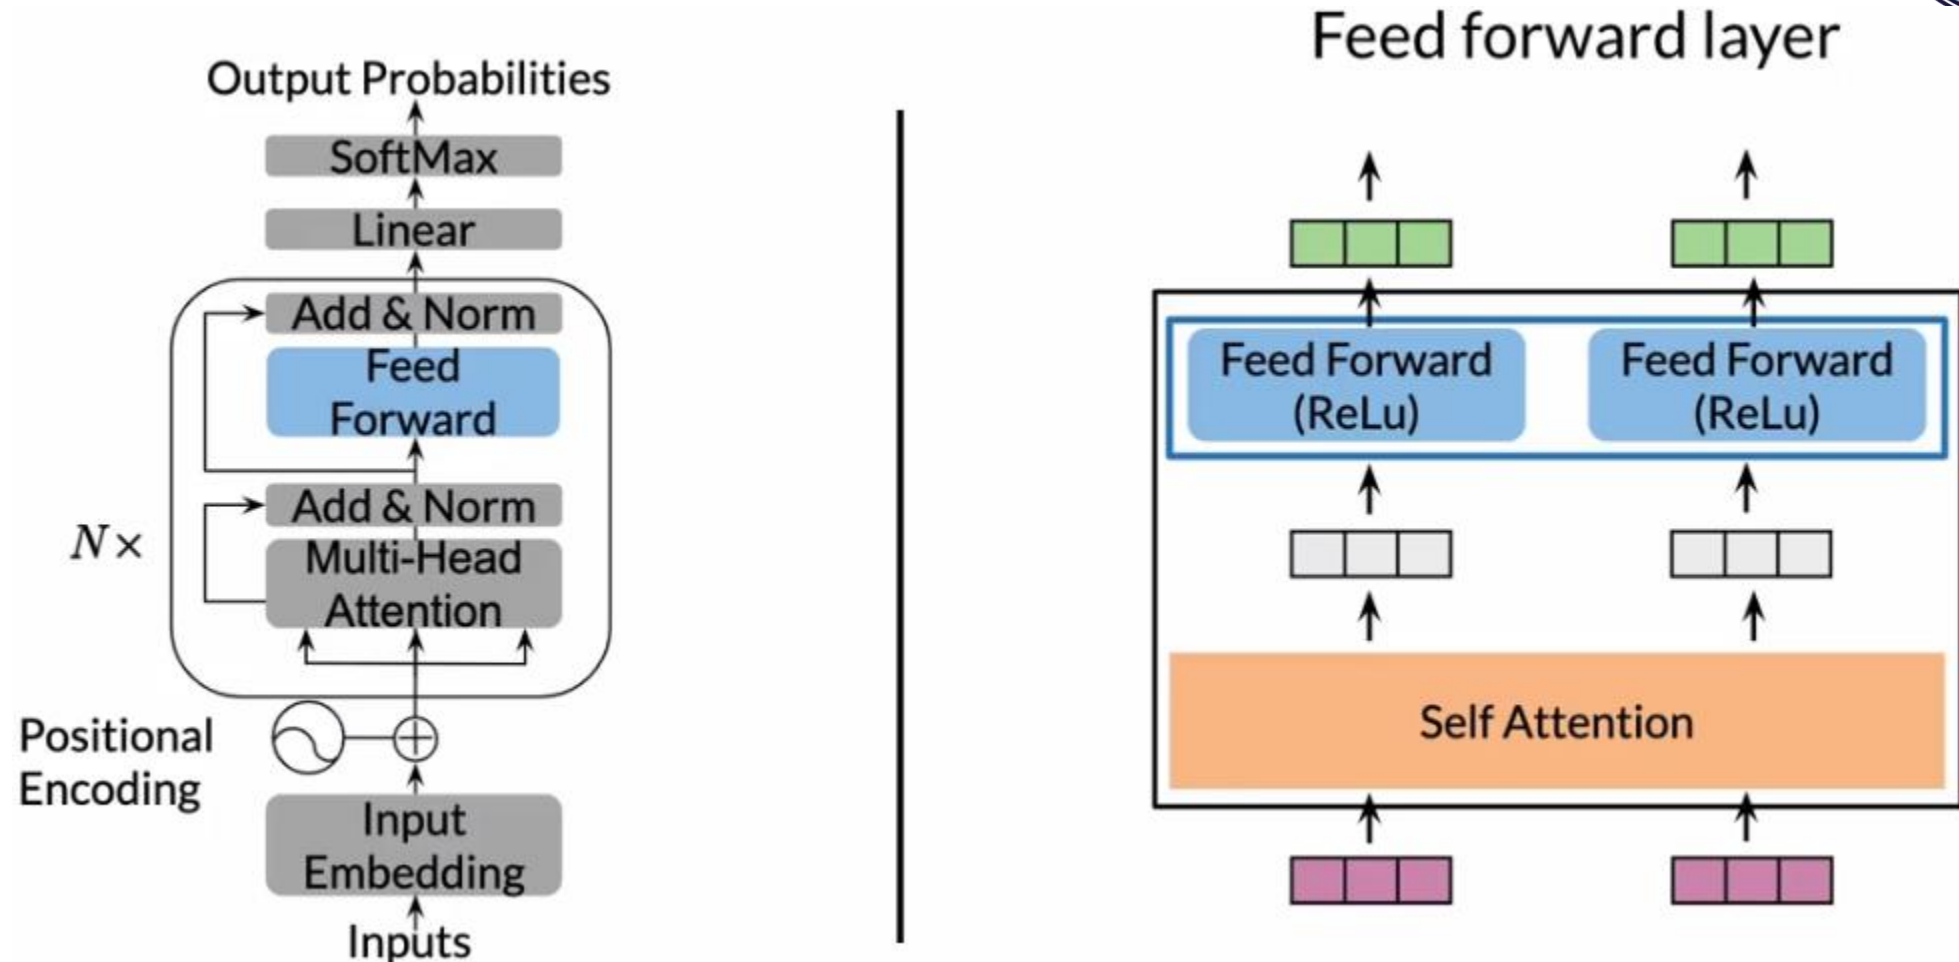


Figure from deeplearning.ai





# Transformer for summarization


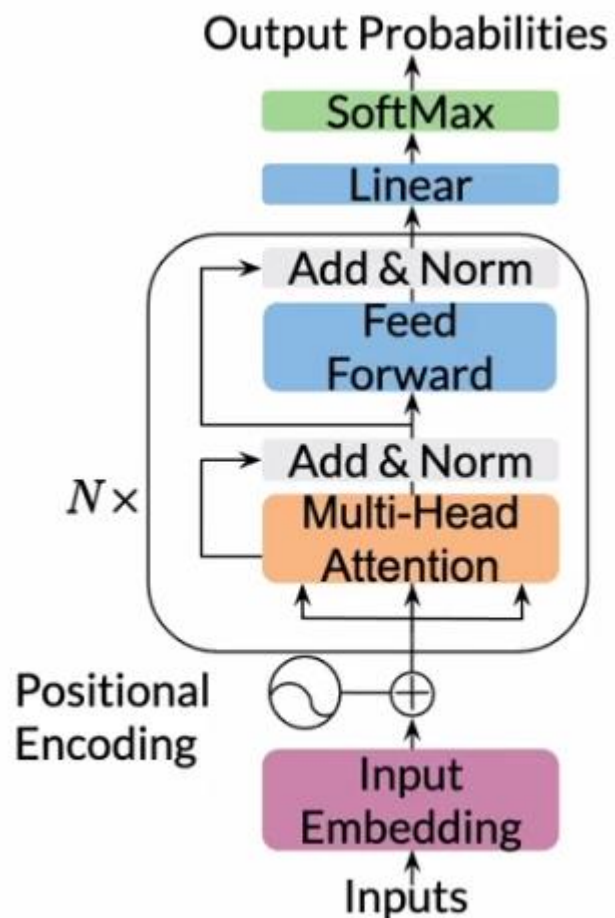
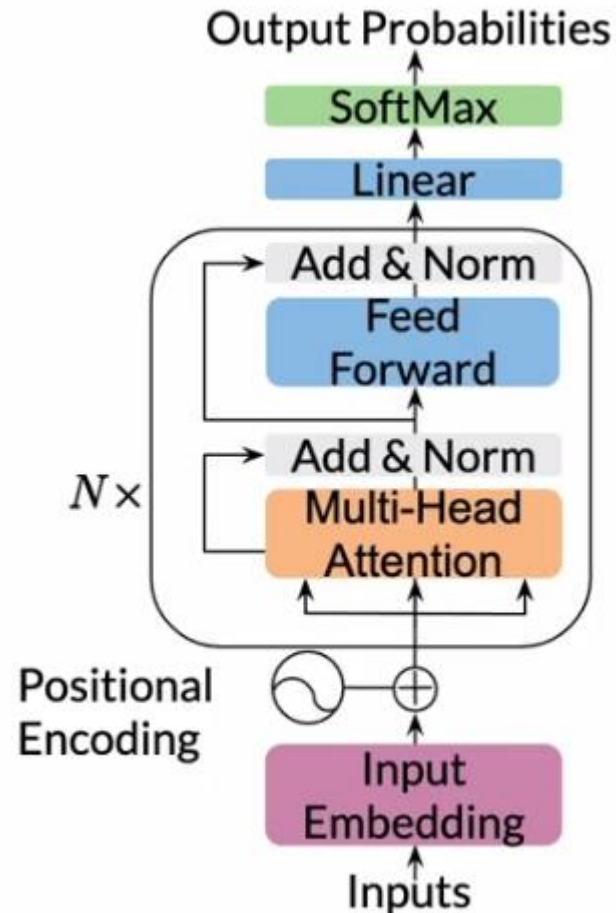


Figure from deeplearning.ai



# Technical details for data processing



## Model Input:

ARTICLE TEXT <EOS> SUMMARY <EOS> <pad> ...

## Tokenized version:

[2, 3, 5, 2, 1, 3, 4, 7, 8, 2, 5, 1, 2, 3, 6, 2, 1, 0, 0]

Loss weights: 0s until the first <EOS> and then 1 on the start of the summary.

when there is little data for the summaries, it actually helps to weight the article loss with non-zero numbers, say 0.2 or 0.5 or even 1.